Importing data into pandas

February 6, 2021

1 Importing data into pandas

There are tons of ways you can get data into a pandas dataframe. Here are a few of the more common ones.

First, let's import pandas `as` pd.

[22]:
```
import pandas as pd
```

1.0.1 From a CSV file

If your data file is delimited with something other than a comma, you'll need to specify that in the `sep` argument. For example, if you had a pipe-delimited file: `pd.read_csv('../data/my-delimited-file.txt', sep='|')`

Let's read in the MLB salary data.

[23]:
```
df_csv = pd.read_csv('../data/mlb.csv')
```

[24]:
```
df_csv.head()
```

[24]:

	NAME	TEAM	POS	SALARY	START_YEAR	END_YEAR	YEARS
0	Clayton Kershaw	LAD	SP	33000000	2014	2020	7
1	Zack Greinke	ARI	SP	31876966	2016	2021	6
2	David Price	BOS	SP	30000000	2016	2022	7
3	Miguel Cabrera	DET	1B	28000000	2014	2023	10
4	Justin Verlander	DET	SP	28000000	2013	2019	7

1.0.2 From a CSV file on the Internet

Just pass in the URL. This example uses licensed child care facility data from Colorado's open data portal.

The values that get returned aren't live – like, if the results changed, your data frame would not update with new values. It reads in the data once.

[25]:
```
df_csv_internet = pd.read_csv('https://data.colorado.gov/api/views/a9rr-k8mu/
↪rows.csv?accessType=DOWNLOAD')
```

[26]:
```
df_csv_internet.head()
```

[26]:

	PROVIDER ID	PROVIDER NAME	PROVIDER SERVICE TYPE	STREET ADDRESS	CITY	STATE	ZIP	COUNTY	COMMUNITY	ECC	…
0	35597.0	Rene Willard	Experienced Family Child Care Home	1473 Walnut St	Windsor	CO	80550	Weld	NaN	Promises for Children	…
1	1670733.0	AUGUSTINE CLASSICAL PRESCHOOL	Child Care Center	480 S Kipling ST	Lakewood	CO	80226	Jefferson	NaN	Triad Early Childhood Council	…
2	1501661.0	DENVER BOTANIC GARDENS SUMMER CAMPS	School-Age Child Care Center	1005 York ST	Denver	CO	80206	Denver	NaN	Denver's Early Childhood Council	…
3	1685302.0	AMANDA DUNCAN	Family Child Care Home	7131 W 75th PL	Arvada	CO	80003	Jefferson	NaN	Triad Early Childhood Council	…
4	4855.0	CCSD DRY CREEK ECS	School-Age Child Care Center	7686 E Hindsdale AVE	Englewood	CO	80112	Arapahoe	NaN	Arapahoe County Early Childhood Council	…

	CCCAP TOTAL AUTH_D1	CCCAP FA STATUS_D1	CCCAP AMOUNT PAID_D1	CCCAP FA EXP DATE_D2	CCCAP TOTAL AUTH_D2	CCCAP FA STATUS_D2	LICENSE FEE DISCOUNT	LONG-LAT	OPERATING STATUS (Self-Report)
0	NaN	False	NaN	NaN	NaN	False	NaN	NaN	Open
1	NaN	False	NaN	NaN	NaN	False	NaN	NaN	Open
2	NaN	False	NaN	NaN	NaN	False	NaN	NaN	NaN
3	NaN	True	NaN	06/30/2022	NaN	True	NaN	NaN	NaN
4	NaN	True	NaN	06/30/2022	NaN	True	NaN	NaN	Open

```
   OPERATING STATUS REPORT DATE
0                    2020-07-31
1                    2020-08-11
2                           NaN
3                           NaN
4                    2020-06-30

[5 rows x 29 columns]
```

1.0.3 From an Excel file

To read an Excel file in pandas, use the `read_excel()` method. Depending on the filetype (`xls` or `xlsx`), you'd also need to separately install into your virtual environment the `xlrd` or `openpyxl` modules. (We've already installed both here.)

You might also want to specify the `sheet_name` to select your worksheet of interest – the default is "the first one."

Here, we're reading in a spreadsheet with data on accidental drug overdoses in Connecticut.

[27]:
```
df_xl = pd.read_excel('../data/CT_Overdoses_2012-2016.xlsx',␣
↪sheet_name='Accidental_Drug_Related_Deaths_')
```

[28]:
```
df_xl.head()
```

[28]:
```
  CaseNumber       Date     Sex   Race   Age Residence City Residence State  \
0   13-16336 2013-11-09  Female  White  53.0         GROTON              NaN
1   12-18447 2012-12-29    Male  White  30.0        WOLCOTT              NaN
2    14-2758 2014-02-18    Male  White  43.0        ENFIELD              NaN
3   14-13497 2014-09-07  Female  White  24.0    WALLINGFORD              NaN
4   13-14421 2013-10-04  Female  White  26.0     WEST HAVEN              NaN

  Residence County   Death City Death State  … Benzodiazepine Methadone  \
0       NEW LONDON       GROTON         NaN  …               Y       NaN
1        NEW HAVEN    WATERBURY         NaN  …             NaN       NaN
2              NaN      ENFIELD         NaN  …               Y       NaN
3              NaN  WALLINGFORD         NaN  …             NaN       NaN
4        NEW HAVEN   WEST HAVEN         NaN  …             NaN       NaN

  Amphet Tramad Morphine (not heroin) Other Any Opioid MannerofDeath  \
0    NaN    NaN                   NaN   NaN        NaN      Accident
1    NaN    NaN                   NaN   NaN        NaN      Accident
2    NaN    NaN                   NaN   NaN        NaN      Accident
3    NaN    NaN                   NaN   NaN        NaN      Accident
4    NaN    NaN                   NaN   NaN        NaN      Accident

  AmendedMannerofDeath                DeathLoc
```

```
0                   NaN   (41.343693, -72.07877)
1                   NaN  (41.554261, -73.043069)
2                   NaN  (41.976501, -72.591985)
3                   NaN  (41.454408, -72.818414)
4                   NaN  (41.272336, -72.949817)

[5 rows x 32 columns]
```

1.0.4 From a Python data collection

Maybe the work you're doing in pandas happens downstream of some other Python processing, so the data exists as a native Python data collection – say, a list of dictionaries. You can turn this (and other Python data collections, like a list of lists) into a pandas dataframe, too.

[29]:
```
test_data = [
    {'name': 'Cody Winchester', 'job': 'Training director', 'location':␣
↪'Colorado Springs, CO'},
    {'name': 'Guy Fieri', 'job': 'Gourmand', 'location': 'Flavortown'},
    {'name': 'Michael Bennet', 'job': 'Senator', 'location': 'Washington, D.C.'}
]
```

[30]:
```
df_py_lod = pd.DataFrame(data=test_data)
```

[31]:
```
df_py_lod.head()
```

[31]:
```
              name                job              location
0  Cody Winchester  Training director  Colorado Springs, CO
1        Guy Fieri           Gourmand            Flavortown
2   Michael Bennet            Senator      Washington, D.C.
```

If you have a list of lists, you would need to also specify the `columns` keyword argument, as well:

[32]:
```
test_data_ls = [
    ['Cody Winchester', 'Training director', 'Colorado Springs, CO'],
    ['Guy Fieri', 'Gourmand', 'Flavortown'],
    ['Michael Bennet', 'Senator', 'Washington, D.C']
]
```

[33]:
```
df_py_lol = pd.DataFrame(data=test_data_ls, columns=['name', 'job', 'location'])
```

[34]:
```
df_py_lol.head()
```

[34]:
```
              name                job              location
0  Cody Winchester  Training director  Colorado Springs, CO
1        Guy Fieri           Gourmand            Flavortown
2   Michael Bennet            Senator       Washington, D.C
```

1.0.5 From an HTML table

OK SO.

This one requires you to install and specify the Python package that has the HTML parsing engine of your choice – BeautifulSoup or lxml. The default is `lxml`, but here we're going to use BeautifulSoup.

Huge caveat! Pulling data directly from an HTML table can be hit and miss, depending on how hairy the underlying HTML is. And if you want to scrape data from a website, it's usually better practice to save the results to a local file, *then* load it up for analysis. But it's good to know that it's an option.

In this example, we've installed `BeautifulSoup` (alias `bs4`) and we're going to import a table of media witnesses to Texas death row executions.

We're going to pass four things to the pandas `read_html()` method: 1. The URL we want to scrape (in quotes, as a string) 2. The `flavor` of parser that we'd like to use to process the HTML (`bs4`) 3. The HTML attributes of the table we're targeting (in this case, the table has a `class` called `tdcj_table`) 4. The number of the list, in the list of lists that gets returned in a dataframe, that has the `header`? (Usually it's 0 – the first one)

Reading through the documentation for this method, we also notice that this method returns a *list* of matching tables as dataframes, so we need to grab the *first* item in this list of tables returned. Our arguments were specific enough that there's only one item in the returned list, though, so we can just grab the first item with `[0]`.

[35]:
```
html_df = pd.read_html('http://www.tdcj.state.tx.us/death_row/
↪dr_media_witness_list.html',
                       flavor='bs4',
                       attrs={'class': 'tdcj_table'},
                       header=0)[0]
```

[36]:
```
html_df.head()
```

[36]:

	Execution	Link	Last Name	First Name	TDCJ Number	Date	\
0	570	Inmate Information	Wardlow	Billy	999137	7/8/2020	
1	569	Inmate Information	Ochoa	Abel	999450	2/6/2020	
2	568	Inmate Information	Gardner	John	999516	1/15/2020	
3	567	Inmate Information	Runnels	Travis	999505	12/11/2019	
4	566	Inmate Information	Hall	Justen	999497	11/6/2019	

	Media Witness List
0	Michael Graczyk, Associated Press; Joseph Brow…
1	Michael Graczyk, Associated Press; Joseph Brow…
2	Michael Graczyk, Associated Press; Joseph Brow…
3	Michael Graczyk, Associated Press; Joseph Brow…
4	Michael Graczyk, Associated Press; Joseph Brow…

1.0.6 From JSON

JSON stands for JavaScript Object Notation. It's a common data interchange format on the web. The `read_json()` method can pull JSON into a data frame.

Pandas can slurp in data from a local `.json` file, or from a URL – say, a JSON API with data on dogs and cats registered in the Sunshine Coast Region of Australia. That one sounds fun! Let's do that.

[37]:
```
json_df = pd.read_json('https://data.sunshinecoast.qld.gov.au/resource/
↪44qj-t4fr.json')
```

[38]:
```
json_df.head()
```

[38]:

	animaltype	name	specificbreed	primarybreed	primarycolour	\
0	D	UNKNOWN DOG NAME	POODLETOY	POODLE	Grey	
1	D	Emie	FOXTERRMN	FOXTER	BlackWhite	
2	D	Jess	MALTESE	MALTES	White	
3	D	Style	BICHONFRS	BICHON	White	
4	D	Benny	SPANIELCV	SPANIE	BlackWhite	

	de_sexed	gender	age	locality
0	Y	F	32	NINDERRY
1	Y	F	21	BURNSIDE
2	Y	M	17	BLI BLI
3	Y	F	23	SIPPY DOWNS
4	Y	M	12	MAROOCHYDORE